


Agent Workflow for Skill-Based Video Generation Acceleration


1 Acceleration Skills




Diffusion
Cache




Token
Pruning



Sparse
Attention



Kernel
Fusion



Quantization

2 Parallel Applicability Executors *



✓ Optimal Implementation

- TeaCache fits best
- Start at 5th step



✗ Not applicable

- Token pruning causes Non-negligible loss on the model



✓ Optimal Implementation

- at density of 10%
- 10th ~ 23rd layers



✓ Optimal Implementation

- Gelu as epilogue
- Fuse repo operation



✓ Optimal Implementation

- Select NVFP4 for B200
- Only on FFN layers

3 Integrator Agent

Selected Skills (Applicable)



Cache



Sparse
Attention



Kernel
Fusion



Quantize



Integration Engine



compose selected skills



global parameter tuning



e2e quality checking



implement stack patch

Validate
&
Iterate



Proposed Acceleration Stack
(Deployment Config + Patches)

4 Human Validator



Validation Dashboard

Efficiency



- latency
- throughput
- memory

Quality



- consistency
- visual fidelity
- alignment



Accept Stack



Revise Policies

*The applicability of certain methods may vary across different models. The specific implementations are provided for illustrative purposes.